

A short description of the model, methods, training configuration used to achieve the final result.

I used Stable Diffusion XL (<https://huggingface.co/stabilityai/stable-diffusion-xl-base-1.0>) as the base model and fine-tuned it using the DreamBooth via LoRA technique (<https://huggingface.co/papers/2106.09685>).

The training dataset comprised all images from the provided set. The training was conducted on GPU A-100, and the network underwent 2400 training steps (480 epochs). Additionally, I utilized mixed precision in FP16 for the training.

Here you can find more info about my model:

https://huggingface.co/maxkochanoff/sdXL_lora_for_portraits

Generated samples:

a girl:



a boy:



a girl and a boy:



a girl with dark hair:



a man with a dog:

